

Concolic Testing of Concurrent Software in the Context of Weak Memory Models

Martin Dobiasch, 0828302

March 25, 2014

1 Motivation & Problem Statement

Software is becoming more and more ubiquitous in our every day life and moreover, the complexity of software is increasing. This poses a major challenge for the reliability of software. To mitigate this problem, research on software testing and verification tries to create automated solutions which exhaustively search for software bugs. Although there has been progress in developing automated testing tools for sequential testing [5], testing concurrent software still remains as a challenge. Concurrency adds an exponential amount of possibilities for tests by adding the scheduling of statements from different threads as an additional factor, i.e., it is also necessary to investigate in which order program statements are executed across different threads.

Weak Memory Models. Most of the test and verification tools for concurrent software rely on the assumption of *sequential consistency*. This means that the execution of a concurrent program is the same as if the statements from different threads had been brought to a sequential order and executed sequentially while the thread-local order remains unchanged for all threads. However, unlike sequential consistency, so-called *Weak Memory Models* (WMMs) allow certain deviations from a thread-local program order. For example, modern processors are allowed to commit the effect of a write operation later and not immediately after the write operation was executed. Commit in this context refers to the fact that the write to a shared variable might be stored in a local buffer before it is stored in global memory. Thus, the executing thread can see the new value of this variable, while other threads might still see the old value. Only when the write is committed then all threads will see the new value. Another interesting aspect

of modern CPUs is that they have multiple pipelines for commands to be executed. This stems from the fact that they have different processing units (e.g. for integer operations, for floating point operations, ...). Whenever an instruction is being processed but has not been finished yet, it is referred to as being in an *in-flight* state. As a result of WMMs, an instruction can remain in an in-flight state while other subsequent instructions have been committed already. By doing so the processor can save time since it does not have to wait for an instruction - like a write - to finish until the next instruction can be executed. This can result in a behaviour that for thread A it seems that thread B does not respect the thread-local order. One aspect that increases the difficulty of understanding these models of computation is that they are usually not published in a comprehensible manner but instead hidden in long technical documentations of a specific processor [6]. One major problem during testing is that testing a software on a machine with a WMM can lead to unexpected program results which are difficult to explain and are hard to reproduce since WMMs allow WMM-specific effects, but do not require the computation to perform them again when running the program again. In summary, WMMs can make identifying faults via manual code inspection extremely hard if not impossible.

Figure 1 illustrates an example with two threads both executing two writes on two shared variables x and y and on two local variables $r1$ and $r2$. Moreover, it defines a state which should usually be infeasible when considering sequential consistency (denoted as 'forbidden'). In the initial state both x and y are initialised with the value 0. Under sequential consistency the forbidden state cannot be reached. However, this state is feasible for ARM and POWER architectures [6]. One possibility to reach this state is as follows: First Thread 0 starts executing the write $x = 1$ but does not commit it. Then the second write $y = 1$ is executed but this time the write is committed immediately. Next, Thread 1 takes over reading the value of y (1) and writing it to $r1$. After this, it executes the second read, thus storing the initial value of x (0) to $r2$. Thereby, the forbidden state is reached.

Thread 0	Thread 1
$x = 1$	$r1 = y$
$y = 1$	$r2 = x$
Initial state: $x = 0 \wedge y = 0$	
Forbidden: $r1 = 1 \wedge r2 = 0$	

Figure 1: Message Passing Example from [6]

Concolic Testing. In past years various testing tools have been developed featuring concolic testing [5], a testing approach that combines the concrete execution of a program with a simultaneous symbolic execution. However, only a few of these tools [9] have the ability to test concurrent software. A testing approach derived from concolic testing which is able to test concurrent software is $(con)^2colic$ testing (*concrete* and *symbolic* execution of a *concurrent* program) [4]. Figure 2 illustrates the high-level architecture of $(con)^2colic$ testing which can mainly be divided into two parts: an execution engine and a reasoning engine. The execution engine gathers symbolic constraints which the reasoning engine then modifies to generate new execution paths. The reasoning engine explores interference scenarios by combining and modifying gathered constraint systems. The $(con)^2colic$ algorithm works in the following way: First, the program under test is executed and (symbolic) information is collected. Based on gathered information alternate program runs are then generated and tested for their realisability. However, there could be an intractable amount of such alternate program runs. In order to reduce the amount of tests, $(con)^2colic$ uses so-called interference scenarios: An *interference scenario (IS)* represents a situation where a program thread needs to read a variable which is written by another thread. For a given IS it can be determined whether a schedule for the threads in the program exists so that this scenario can be realised. Thus, the number of schedules to be tested can be reduced dramatically, while the testing routine is still complete (under some standard concolic testing assumptions). The tool CONCREST [4] implements $(con)^2colic$ testing as an extension of CREST¹. However, CONCREST has no means to test programs with respect to the effects of weak memory models and instead relies on sequential consistency. As a consequence of this, certain aspects of concurrent behaviour remain untested. *The problem to be solved in this thesis is to add support for weak memory models in $(con)^2colic$ testing.*

2 Expected Results

In this thesis I will extend the execution platform of CONCREST with the capability to simulate effects of weak memory models. It will provide an interface to enable the precise control of the effects of a WMM.

Since there exist different WMMs, the extension has to be configurable with regard to which WMM should be used while performing program executions. This also involves identifying important/relevant WMMs and in-

¹<http://code.google.com/p/crest/>

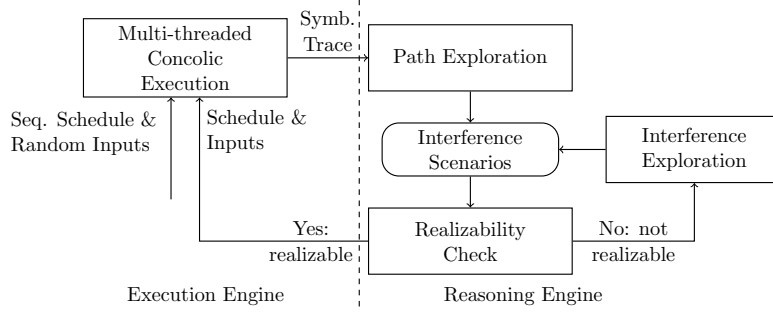


Figure 2: **Overview of (Con)²colic Testing [4].**

vestigating their properties.

A major feature of CONCREST is that during the performed tests it is possible to use dynamic analysis tools, like Valgrind, and one goal is to preserve this feature.

3 Methodological Approach

1. **Literature Review:** In a first step I will gather background information about weak memory models and (con)²colic testing by a thorough review of existing literature. During this literature review, I will identify existing WMMs and study their similarities and differences. By doing so, I will identify central concepts of WMMs like reorderings. Later on, it has to be decided which of these models have to be supported.
2. **Model:** In a next step I will define a formalism for modelling WMMs. These models should then be loadable by the testing tool. This will not only ease the development processes of the tool, but will also provide extensibility for additional WMMs in the future.
3. **Implementation:** I will extend CONCREST with a simulation capability which is able to inject effects of a WMM into a concolic execution. In order to enable the reasoning engine to precisely control WMM effects, I will realise a WMM-specific interface between reasoning and execution engine.
4. **Evaluation:** I will evaluate the implementation on a set of benchmark programs and WMMs. The benchmarks will be based on findings during my literature review.

4 State of the Art

There are various other tools for software testing and verification. Like CONCREST, SAGE [5] and CUTE [9] implement the DART (Directed Automated Random Testing) approach but have no support for multi-threading. While the test tool CUTE (Concolic Unit Testing and Explicit Path Model-Checking) can only test sequential C programs, jCUTE can test concurrent Java programs for data races, deadlocks, infinite loops and uncaught exception errors [9]. ConTest [3] takes a different approach: It takes a Java program and a test for the program and tries to increase the coverage of two test models. The first model just requires a context switch for a method in the Java program, while the second model records context switches forced by a specific method in the Java program. ConTest tests a program only on a single core CPU since any concurrency issue on a single execution is also a bug for a multi-core execution. The tool instruments the program using sleep, yield and priority. During the tests the instrumentations are used to force a context switch from one thread to a different thread [3]. Poirot [8] takes a similar approach as ConTest and translates a concurrent program into a sequential execution. This sequential execution is forced to contain context switches which are considered to be crucial for possible bugs. The sequential program is then analysed using static analysis techniques. Chess [7] uses model checking for testing concurrent programs. It tries to enumerate possible thread schedules giving priority to schedules with fewer preemptions. To incorporate Chess into an existing test the code has to be modified slightly. The framework provides wrapper functions for common concurrency API-functions such as the Win32 API such that the execution of the test can be instrumented and recorded. Fusion [11, 10] also records a concurrent trace of a program execution but then transforms this trace into a logical formula that also encodes certain properties to be checked. If the formula has a solution then a property violation has been determined or otherwise the trace does not violate the given properties. As an extension, the concept of interference abstractions, a concept similar to interference scenarios, is introduced [10].

None of the previously mentioned tools is able to test using WMMs. CBMC has an extension which makes bounded model checking with respect to a given WMM possible. One problem with this approach is that bounded model checking does not scale well for larger concurrent software [1]. Another tool which is capable of testing software with respect to WMM properties is RELAXER [2], a tool facilitating an active random testing technique and combining it with the search for data races using different WMMs. This

Tool	Basic principle	WMM support	Source code available
CBMC [1]	bounded model checking	yes	yes
CREST ²	concolic testing	no	no
CONCREST [4]	(con) ² colic testing	no	yes
SAGE [5]	concolic testing	no	no
Cute [9]	concolic testing	no	yes
jCute [9]	concolic testing	no	yes
CHESS [7]	model checking	no	yes
Poirot [8]	static analysis	no	no
Fusion [11, 10]	model checking	no	no
ConTest [3]	dynamic analysis	no	no
RELAXER [2]	dynamic analysis	yes	no

Table 1: Overview of concurrency testing tools

means that the program uses an analysis to predict potential bugs. After this initial analysis the program tries to trigger these bugs by actively controlling the scheduler. Table 1 summarizes the above mentioned tools.

5 Relevance to the Curriculum of Computational Intelligence

This thesis fits into the Computational Intelligence module *Programming Languages and Verification* as it requires deeper knowledge of the *semantics of programming languages* (such as C and C++), *software testing* and *software verification*.

References

- [1] Jade Alglave, Daniel Kroening, and Michael Tautschnig, *Partial orders for efficient bounded model checking of concurrent software*, CAV (Natasha Sharygina and Helmut Veith, eds.), Lecture Notes in Computer Science, vol. 8044, Springer, 2013, pp. 141–157.
- [2] Jacob Burnim, Koushik Sen, and Christos Stergiou, *Testing concurrent programs on relaxed memory models*, ISSTA (Matthew B. Dwyer and Frank Tip, eds.), ACM, 2011, pp. 122–132.

²<http://code.google.com/p/crest/>

- [3] Orit Edelstein, Eitan Farchi, Yarden Nir, Gil Ratsaby, and Shmuel Ur, *Multithreaded java program test generation*, IBM Systems Journal **41** (2002), no. 1, 111–125.
- [4] Azadeh Farzan, Andreas Holzer, Niloofar Razavi, and Helmut Veith, *Con2colic testing*, ESEC/SIGSOFT FSE (Bertrand Meyer, Luciano Baresi, and Mira Mezini, eds.), ACM, 2013, pp. 37–47.
- [5] Patrice Godefroid, Nils Klarlund, and Koushik Sen, *Dart: directed automated random testing*, PLDI (Vivek Sarkar and Mary W. Hall, eds.), ACM, 2005, pp. 213–223.
- [6] Luc Maranget, Susmit Sarkar, and Peter Sewell, *A Tutorial Introduction to the ARM and POWER Relaxed Memory Models*.
- [7] Madanlal Musuvathi, Shaz Qadeer, and Thomas Ball, *CHESS: A Systematic Testing Tool for Concurrent Software*, 2007.
- [8] Shaz Qadeer, *Poirot - a concurrency sleuth*, Formal Methods and Software Engineering - 13th International Conference on Formal Engineering Methods, ICFEM 2011, Durham, UK, October 26-28, 2011. Proceedings (Shengchao Qin and Zongyan Qiu, eds.), Lecture Notes in Computer Science, vol. 6991, Springer, 2011, p. 15.
- [9] Koushik Sen, Darko Marinov, and Gul Agha, *Cute: a concolic unit testing engine for c*, ESEC/SIGSOFT FSE (Michel Wermelinger and Harald Gall, eds.), ACM, 2005, pp. 263–272.
- [10] Nishant Sinha and Chao Wang, *On interference abstractions*, Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages (New York, NY, USA), POPL '11, ACM, 2011, pp. 423–434.
- [11] Chao Wang, Sudipta Kundu, Rhishikesh Limaye, Malay Ganai, and Aarti Gupta, *Symbolic predictive analysis for concurrent programs*, Formal Aspects of Computing **23** (2011), no. 6, 781–805 (English).

(Unterschrift Betreuer)